

ExerciseSheet1

November 7, 2021

1 IPSA 2021 - Exercise 1

The goal of this exercise is to get used to practical image processing in *python / numpy / scipy*.

Do **NOT** use additional third party libraries such as *OpenCV* or *scikit-image* for the coding tasks in this course!

1.1 0 - Warm Up

```
In [1]: import imageio
import numpy as np
import scipy.ndimage as img
import matplotlib.pyplot as plt

import timeit, functools
```

```
In [ ]: !wget https://github.com/FabriceBeaumont/1-MA-INF_2314_IPSA_Repo/blob/main/Exercises/'

--2021-11-05 11:46:07--  https://github.com/FabriceBeaumont/1-MA-INF_2314_IPSA_Repo/blob/main/
Resolving github.com (github.com)... 140.82.113.4
Connecting to github.com (github.com)|140.82.113.4|:443... connected.
HTTP request sent, awaiting response... 404 Not Found
2021-11-05 11:46:07 ERROR 404: Not Found.
```

```
In [2]: def imageRead(imgname, pilmode='L', arrtype=np.float):
        """
        Read an image file into a numpy array

        imgname: str
            name of image file to be read
        pilmode: str
            for luminance / intensity images use L
            for RGB color images use RGB

        arrtype: numpy dtype
            use np.float, np.uint8, ...
        """
```

```

return imageio.imread(imgname, pilmode=pilmode).astype(arrtype)

def imageWrite(arrF, imgname, arrtype=np.uint8):
    """
    Write a numpy array as an image file
    the file type is inferred from the suffix of parameter imgname, e.g. .png
    arrF: array_like
        array to be written
    imgname: str
        name of image file to be written
    arrtype: numpy dtype
        use np.uint8, ...
    """
    imageio.imwrite(imgname, arrF.astype(arrtype))

```

To display an intensity image use:

```

In [ ]: arrF = imageRead("portrait.png")
        plt.imshow(arrF / 255, cmap='gray')
        plt.xticks([]); plt.yticks([])
        plt.show()

```



To display an (RGB) color image use:

```

In [ ]: arrF = imageRead('asterixRGB.png', pilmode='RGB')
        plt.imshow(arrF / 255)
        plt.xticks([]); plt.yticks([])
        plt.show()

```



1.2 1 - The emboss effect

```
In [48]: filename = "portrait.png"
```

```
arrF = imageRead(filename)
print(f"The image '{filename}' has {arrF.shape[0]} rows and {arrF.shape[1]} columns.")
```

The image 'portrait.png' has 256 rows and 256 columns.

1.2.1 Different emboss implementations

```
In [49]: def embossV1(arrF):
```

```
    """
```

```
    Naive Emboss: Iterate through all cells and explicitly compute the
    every gradient. Do so by computing the difference between the intensity
    values in the lower right and upper left cell.
    Add 255/2 = 128 to compensate for the induced value shift.
    Clip the results to the interval [0, 255].
```

```
    Slow implementation.
```

```
    """
```

```
    M, N = arrF.shape
    arrG = np.zeros((M,N))
```

```
    for i in range(1,M-1):
        for j in range(1,N-1):
```

```

arrG[i,j] = 128 + arrF[i+1,j+1] - arrF[i-1,j-1]
arrG[i,j] = np.maximum(0, np.minimum(255, arrG[i,j]))

```

```

return arrG

```

```

In [50]: def embossV2(arrF):

```

```

    """

```

```

    Numpythonic Emboss: Use python-slicing to compute all gradients
    at once by subtracting the upper left (M-2 x N-2)-matrix
    from the lower right (M-2 x N-2)-matrix - to yield the gradients
    of the central (M-1 x N-1)-matrix.

```

```

    Add 255/2 = 128 to compensate for the induced value shift.

```

```

    Clip the results to the interval [0, 255].

```

```

    Fast implementation.

```

```

    """

```

```

    M, N = arrF.shape

```

```

    arrG = np.zeros((M,N))

```

```

    arrG[1:M-1,1:N-1] = 128 + arrF[2:,2:] - arrF[:-2,:-2]

```

```

    arrG = np.maximum(0, np.minimum(255, arrG))

```

```

    return arrG

```

```

In [51]: def embossV3(arrF):

```

```

    """

```

```

    Convolution Emboss: similar to CNNs, use a 3x3-mask which sweeps
    over all non-border-cells.

```

```

    Add 255/2 = 128 to compensate for the induced value shift.

```

```

    Clip the results to the interval [0, 255].

```

```

    Fast and practical (generalizable) implementation.

```

```

    """

```

```

    mask = np.array([[ -1, 0,  0],
                     [ 0, 0,  0],
                     [ 0, 0, +1]])

```

```

    arrG = 128 + img.correlate(arrF, mask, mode='reflect')

```

```

    arrG = np.maximum(0, np.minimum(255, arrG))

```

```

    return arrG

```

```

In [69]: def embossV4(arrF):

```

```

    """

```

```

    Numpythonic Emboss: Use python-slicing to compute all gradients
    at once by subtracting the upper left (M-1 x N-1)-matrix
    from the lower right (M-1 x N-1)-matrix - to yield the gradients
    of the central (M-1 x N-1)-matrix. Notice, that only this matrix is stored,

```

```

    of the central (M-1 x N-1)-matrix. Notice, that only this matrix is stored,

```

thus the dimensions are diminished in each direction by two and no black bordered will remain.

Add $255/2 = 128$ to compensate for the induced value shift.

*Clip the results to the interval $[0, 255]$, by using
python's array boolean indexing. (Masks for the corner values.)
"""*

```
arrG = 128 + arrF[2:,2:] - arrF[:-2,:-2]
arrG[arrG< 0] = 0
arrG[arrG>255] = 255
```

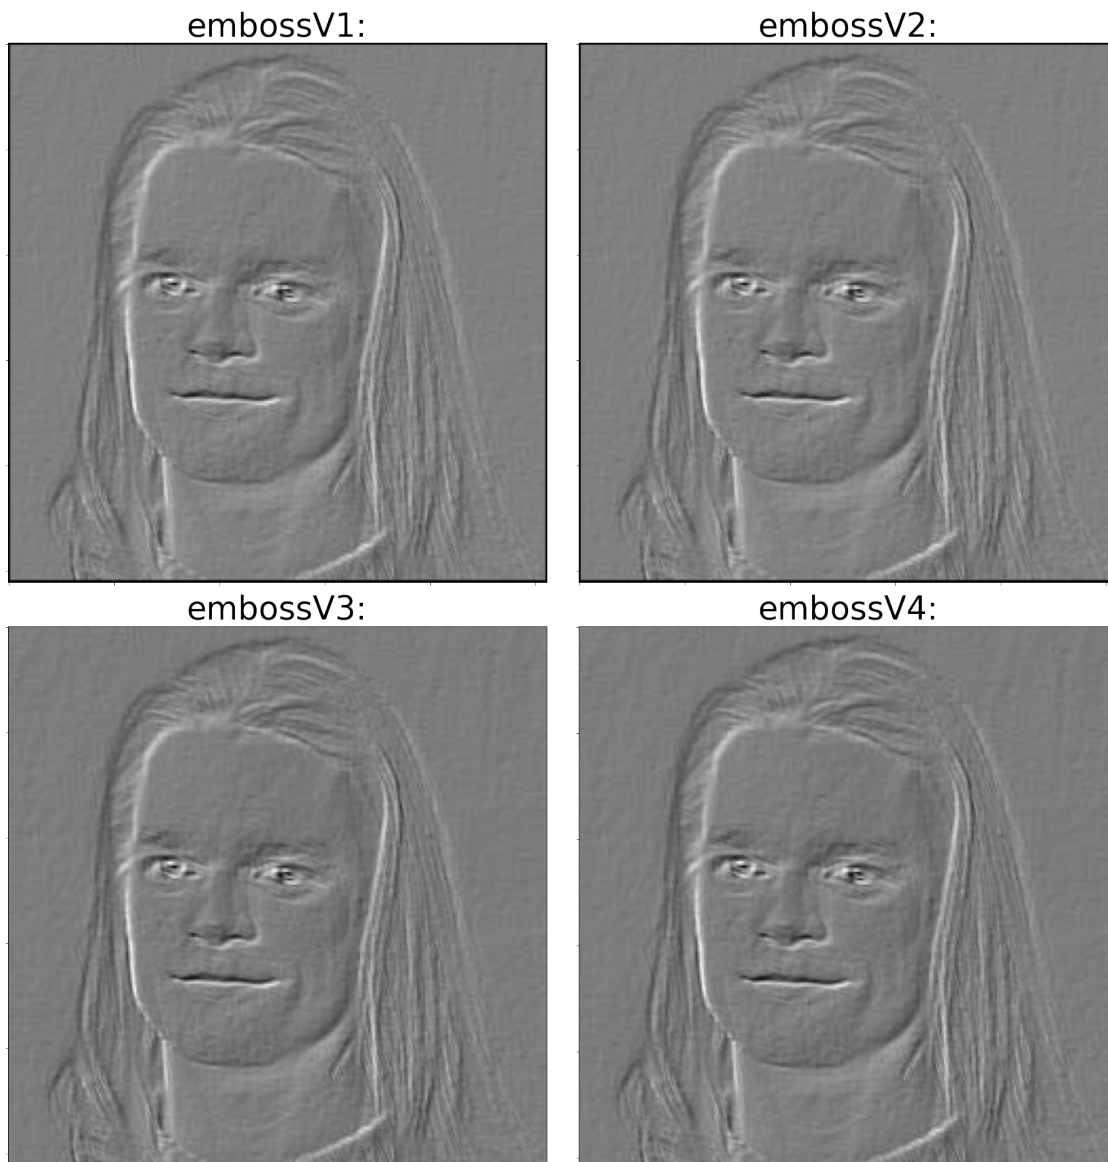
```
return arrG
```

```
In [70]: methods = [embossV1, embossV2, embossV3, embossV4]
arrGs = [arrF] * len(methods)
```

```
fig, axs = plt.subplots(int(0.5*len(methods)), 2, figsize=(20,20))
```

```
for i, mtd in enumerate(methods):
    arrGs[i] = mtd(arrF)
```

```
ax = axs[int(i/2), i%2]
ax.imshow(arrGs[i], cmap='gray')
ax.set_title(f"{mtd.__name__}:", fontsize=40)
ax.set_xticklabels([]); ax.set_yticklabels([])
fig.tight_layout()
```



Notice that method embossV3 and embossV4 do not result in a black border.
Furthermore notice, how embossV3 does not trim the image:

```
In [68]: for i, arrG in enumerate(arrGs):
          print(f"Image dimensions when using method embossV{i}: {arrG.shape}")
          print("\tDoes %screate a black border.\n\n"%( "not " if arrG[0][100] != 0 else ""))
```

```
Image dimensions when using method embossV0: (256, 256)
\tDoes not create a black border.
```

```
Image dimensions when using method embossV1: (256, 256)
\tDoes create a black border.
```

Image dimensions when using method embossV2: (256, 256)
Does not create a black border.

Image dimensions when using method embossV3: (254, 254)
Does not create a black border.

1.3 2 - Timing the emboss effect

1.3.1 2 a)

```
In [ ]: def runtime_evaluation(data, methods, nRep, nRun):
        print("Mean runtimes for:")
        for method in methods:
            ts = timeit.Timer(functools.partial(method, data)).repeat(nRep, nRun)
            print(f"{method.__name__}: \t{min(ts) / nRun} \t[sec]")

In [ ]: mtds = [embossV1, embossV2, embossV3, embossV4]
        nRep = 3
        nRun = 100

In [ ]: filename0 = "portrait.png"
        arrF0 = imageRead(filename0)

        runtime_evaluation(arrF0, methods, nRep, nRun)
```

Mean runtimes for:

embossV1:	0.29045785174999766	[sec]
embossV2:	0.0006378255400022681	[sec]
embossV3:	0.000883353970002645	[sec]
embossV4:	0.00031468285000300965	[sec]

1.3.2 2 b)

```
In [ ]: filename1 = "asterix.png"
        arrF1 = imageRead(filename1)

        runtime_evaluation(arrF1, methods, nRep, nRun)
```

Mean runtimes for:

embossV1:	1.646454750269995	[sec]
embossV2:	0.0034446558699983143	[sec]
embossV3:	0.004903635989994655	[sec]
embossV4:	0.0020958068999971146	[sec]

1.3.3 2 c)

```
In [ ]: n0, m0 = arrF0.shape
        n1, m1 = arrF1.shape
        print(f"The image '{filename0}' has dimensions {n0}x{m0} and thus {n0*m0} cells.")
        print(f"The image '{filename1}' has dimensions {n1}x{m1} and thus {n1*m1} cells.")
        print(f"That is almost {round(n1*m1/(n0*m0))}-times as much!")
        print(f"This corresponds to the increase in runtime, approximately linear wrt to the n
```

The image 'portrait.png' has dimensions 256x256 and thus 65536 cells.

The image 'asterix.png' has dimensions 496x730 and thus 362080 cells.

That is almost 6-times as much!

This corresponds to the increase in runtime, approximately linear wrt to the number of cells.

1.4 3 - Working with RGB color images

```
In [ ]: arrF = imageRead('asterixRGB.png', pilmode='RGB')

        print(arrF.shape)
        plt.imshow(arrF / 255)
        plt.xticks([]); plt.yticks([])
        plt.show()
```

(496, 730, 3)




```
In [ ]: arrG = np.copy(arrF)
        arrG[:, :, 0] = 0

        print(arrG.shape)
        plt.imshow(arrG / 255)
        plt.xticks([]); plt.yticks([])
        plt.show()

(496, 730, 3)
```



We can see, that all values of the third RGB-channel were set to zero. Thus the **image does no longer contain the color red**. Red is replaced by black.

Notice that asterix's trousers are black. And the white background is only composed out of the colors green and blue.

Explanation of the code `arrG[:, :, 0] = 0` operates on all cells (`[:, :]`) and the first (zeroth) of the color channels (which are stored in the third position). Thus the red-value for every pixel is set to zero/black.


1.5 4 - Getting used to slicing (part 1)

```
In [ ]: arrF = imageRead('asterix.png')
        arrG = imageRead('portrait.png')

        arrH = np.copy(arrF)
        n, m = arrG.shape

        i, j = 100, 200
        arrH[i:i+n, j:j+m] = arrG
```


```
In [ ]: plt.imshow(arrH / 255, cmap='gray')
plt.xticks([]); plt.yticks([])
plt.show()
```



1.6 5 - Getting used to slicing (part 2)

```
In [ ]: arrF = imageRead('portrait.png')
n, m = arrF.shape
d = 16
arrF[:, ::d] = 0
arrF[:, :d, :] = 0

In [ ]: plt.imshow(arrF / 255, cmap='gray')
plt.xticks([]); plt.yticks([])
plt.show()
```



1.7 6 - Getting used to meshgrids (part 1)

```
In [ ]: arrF = imageRead('portrait.png')
        n, m = arrF.shape
        width_offset = 50
        height_offset = 85
        ci, cj = 128, 110

In [ ]: # Use meshgrid to create an index mask.
        rs, cs = np.meshgrid(np.arange(m), np.arange(n), indexing='ij')

        # For the ellipse use a similar formula like the unit circle
        # Simply offset the origin by the defined center, and
        # scale the height (y-axis) and width (x-axis) differently.
        mask = ((rs-ci)**2 / height_offset**2 + (cs-cj)**2 / width_offset**2 <= 1)

        # Apply the change to the image - in the area of the mask.
        arrF[mask] = 255

In [ ]: plt.imshow(arrF / 255, cmap='gray')
        plt.xticks([]); plt.yticks([])
        plt.show()
```



1.8 7 - Getting used to meshgrids (part 2)

```
In [45]: arrF = imageRead('portrait.png')
         n, m = arrF.shape
         radii = [32, 16, 64]

In [46]: def circle_cutouts(arrF, r=32):
         m, n = arrF.shape

         # Create a meshgrid for only one circle (with diameter 2*r).
         rs, cs = np.meshgrid(np.arange(2*r), np.arange(2*r), indexing='ij')

         # Define the circle cutout.
         # Notice that the OUTSIDE of the circle is masked ('>').
         mask = ((rs-r)**2 + (cs-r)**2 > r**2)

         # Tile the circle as often as it fits in the dimensions
         mask = np.tile(mask, (m//(2*r), n//(2*r)))

         # Apply the mask
         arrG = np.copy(arrF)
         arrG[mask] = 255

         return arrG

In [ ]: fig, axs = plt.subplots(1, 3, figsize=(20,20))
```


```

for i, radius in enumerate(radII):

    arrG = circle_cutouts(arrF, r=radius)

    ax = axs[i]
    ax.imshow(arrG, cmap='gray')
    ax.set_title(f"r={radius}", fontsize=35)
    ax.set_xticklabels([]); ax.set_yticklabels([])
    fig.tight_layout()

```



Runtime measurements

```

In [47]: nRep = 3
         nRun = 100
         ts = timeit.Timer(func tools.partial(circle_cutouts, arrF)).repeat(nRep, nRun)
         print(f"{min(ts) / nRun} [sec]")

```

0.00025478838999788423 [sec]

```

In [ ]: # TODO: Evaluate runtime complexity

```

1.9 8 - Getting used to meshgrids (part 3)

```

In [12]: m = n = 256

```

```

def sin_func(x, y, nu, n):
    sinus_term = np.sin(2 * np.pi * nu * x / (n-1) + np.pi / 2 * (n-1))
    return 0.5 * (sinus_term + 1) * 255

```

```

In [43]: xs, ys = np.meshgrid(np.arange(n), np.arange(m))


```

```
In [32]: nu_list = [1., 2., 3., 9.]

fig, axs = plt.subplots(1, len(nu_list), figsize=(20,20))

for i, nu in enumerate(nu_list):
    arrF = sin_func(xs, ys, nu, n)

    ax = axs[i]
    ax.imshow(arrF, cmap='gray')
    ax.set_title(r"$\nu$=%i"%nu, fontsize=40)
    ax.set_xticklabels([]); ax.set_yticklabels([])
    fig.tight_layout()
```



1.10 9 - Getting used to meshgrids (part 4)


```
In [31]: def dampedVerticalSineWave(M, N, nu):
    xs, ys = np.meshgrid(np.arange(N), M-1-np.arange(M))
    return 0.5 * (np.exp(-4 / M * ys) * np.sin(2 * np.pi * nu * ys / (M-1) + np.pi / 2))

In [33]: nu_list = [5., 9., 11., 13.]
```

```
fig, axs = plt.subplots(1, len(nu_list), figsize=(20,20))

for i, nu in enumerate(nu_list):
    arrF = dampedVerticalSineWave(256, 256, nu)

    ax = axs[i]
    ax.imshow(arrF, cmap='gray')
    ax.set_title(r"$\nu$=%i"%nu, fontsize=40)
    ax.set_xticklabels([]); ax.set_yticklabels([])
    fig.tight_layout()
```



1.11 10 - Getting used to universal functions

```
In [34]: def solarize(arrF, nu):
          return np.cos(arrF * nu * 2 * np.pi / 255.) * 127.5 + 127.5
```

```
In [42]: nu_list = [0.5, 1.0, 1.5, 2.0]
```

```
fig, axs = plt.subplots(1, len(nu_list), figsize=(20,20))
arrF = imageRead('portrait.png')
```

```
for i, nu in enumerate(nu_list):
    arrG = solarize(arrF, nu)
```

```
    ax = axs[i]
    ax.imshow(arrG, cmap='gray')
```

```
    ax.set_title(r"$\nu$="+ f"{nu:.1f}", fontsize=40)
    ax.set_xticklabels([]); ax.set_yticklabels([])
    fig.tight_layout()
```

